

Qwen2.5-VL-7B zero-shot ViewTree depth-1 (MindCube-trained) ViewTree-D beam (corpus-trained)
SFT-plain (benchmark data) corpus SFT, single pass

