

one reconstruction serves every question and every search path in the scene; prefill dominates VLM cost, so the runtime batches and shortens prefills

