


all labels come from answer outcomes on the training corpus; no human annotation of camera trajectories



+ 0..3
rendered
views



same correct answer supervises
every trajectory prefix

① Answering



try valid movements; keep the
shortest trajectory that fixes
the answer (ground truth known);
correct at start → Stop, depth 0

② Oracle search

camera controller:
imitate oracle actions + Stop

confidence head: predict
answer correctness per state

optional RL over trajectories
(kept as ablation)

③ Policy & confidence