



1 Captured scene



4 of 64 recorded frames: the bed is covered



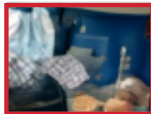
current view



reconstruction (top view)

request 1
turn left $\approx 85^\circ$

World-model
prediction



× invents a bed and walls that
contradict the recorded frames

Actual view
(captured frame)



ViewTree
render

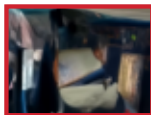


✓ same geometry as the
actual view, sparser texture

same stored geometry
for every request

prediction fed back
as the next input:
errors compound

request 2
then forward
 ≈ 0.7 m



× drifts further from its
own invented scene



✓ no error accumulation